

Advance Operating System

Assignment 2 Report

By: Shlok Bansal (2021201046) Mtech CSE

Introduction:

In this assignment, we have to create 4 system calls for different purposes for linux Kernel **v4.19.210**

In the following report I have attached all the steps that I have followed to achieve this objective

Downloading And Extracting the Kernel

Downloading The kernel

```
wget https://www.kernel.org/pub/linux/kernel/v4.x/linux-4.19.210.tar.xz
```

The above wget command is used to download linux kernel 4.19.210 from kernel.org website using terminal.

Extracting the Kernel

```
sudo tar -xvf linux-4.19.210.tar.xz -C/usr/src/
```

- The tar command is used to extract the tar file which we have downloaded with the help of the above command.
- -x : extract files from an archive.
- -v : requested using the -verbose option, when extracting archives.
- -f : file archive; use archive file or device archive
- -C : extract to the directory specified after it.

The above command will extract the linux-4.19.210.tar.xz file in /usr/src directory. All the kernels are present inside this /usr/src directory. Sudo is written as sudo privileges are required to make any change in this directory.

Note: To avoid writing sudo after every command when we need to make change we are going to go to root via **sudo -s** command.

```
sudo -s
```

Question 1) Write syscall to print welcome message to Linux logs.

Step 1) Go to the kernel directory where we want to make a new system call

```
cd /usr/src/linux-4.19.210/
```

With this we change our and work in the kernel directory.

Step 2) Make the directory in which you will write the source code of the system call and go into that directory.

```
mkdir shlokhello
```

```
cd shlokhello
```

Step 3) Create shlokhello.c in this directory.

```
gedit shlokhello.c
```

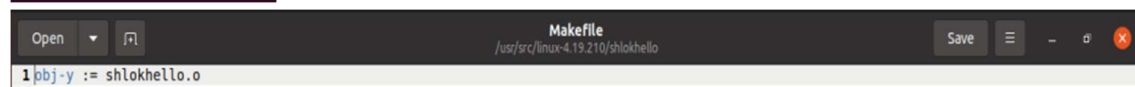


```
1 #include <linux/kernel.h>
2
3 asmlinkage long shlokhello(void)
4 {
5     printk("Welcome to Linux shlok");
6     return 0;
7 }
```

- **asmlinkage** is a keyword which is used to indicate that all parameters of the function would be available of stack.
- **void** indicates function that shlokhello() does not take any parameters.
- **printk** is a function like printf but it is used to print in kernel log.

Step 4) Create Makefile

```
gedit Makefile
```



```
1 obj-y := shlokhello.o
```

This is to ensure that shlokhello.c is compiled when the kernel is compiled.

Step 5) Edit the Makefile of Kernel

```
cd /usr/src/linux-4.19.210/
```

```
nano Makefile
```

Look for the line which looks like below

```
core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/
```

In this line add shlokhello/

```
core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ hello/ shlokhello/
```

- By adding shlokhello/ in this line we are telling the compiler of the kernel that our new system call is defined in shlokhello directory.
- It will tell the compiler to make the Makefile in this directory which was mentioned above.

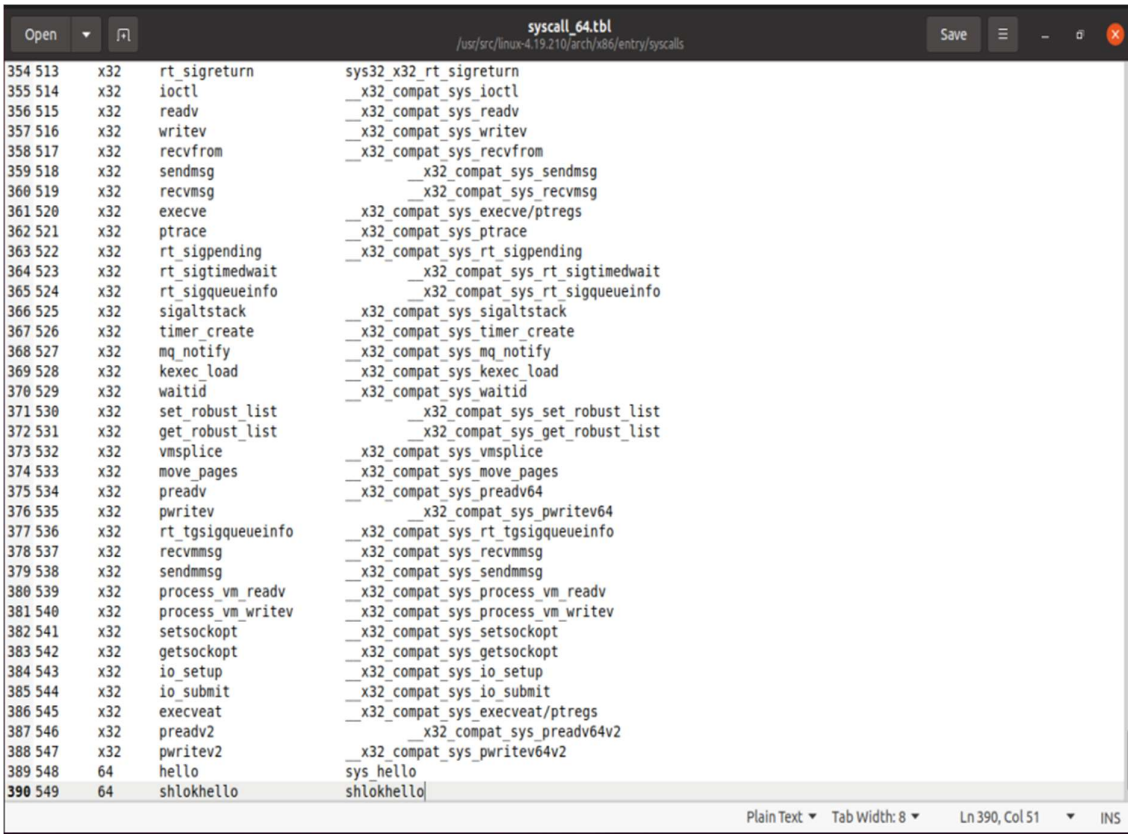
Step 6) Add the new system call to the system call table of kernel

```
cd arch/x86/entry/syscalls/
```

The system call table of the kernel lies in above directory. Change your current directory with the above command.

As my system is 64 bit so add entry in syscall_64.tbl .

```
gedit syscall_64.tbl
```



```
354 513 x32 rt_sigreturn sys32_x32_rt_sigreturn
355 514 x32 ioctl _x32_compat_sys_ioctl
356 515 x32 readv _x32_compat_sys_readv
357 516 x32 writev _x32_compat_sys_writev
358 517 x32 recvfrom _x32_compat_sys_recvfrom
359 518 x32 sendmsg _x32_compat_sys_sendmsg
360 519 x32 recvmsg _x32_compat_sys_recvmsg
361 520 x32 execve _x32_compat_sys_execve/ptregs
362 521 x32 ptrace _x32_compat_sys_ptrace
363 522 x32 rt_sigpending _x32_compat_sys_rt_sigpending
364 523 x32 rt_sigtimedwait _x32_compat_sys_rt_sigtimedwait
365 524 x32 rt_sigqueueinfo _x32_compat_sys_rt_sigqueueinfo
366 525 x32 sigaltstack _x32_compat_sys_sigaltstack
367 526 x32 timer_create _x32_compat_sys_timer_create
368 527 x32 mq_notify _x32_compat_sys_mq_notify
369 528 x32 kexec_load _x32_compat_sys_kexec_load
370 529 x32 waitid _x32_compat_sys_waitid
371 530 x32 set_robust_list _x32_compat_sys_set_robust_list
372 531 x32 get_robust_list _x32_compat_sys_get_robust_list
373 532 x32 vmsplice _x32_compat_sys_vmsplice
374 533 x32 move_pages _x32_compat_sys_move_pages
375 534 x32 preadv _x32_compat_sys_preadv64
376 535 x32 pwritev _x32_compat_sys_pwritev64
377 536 x32 rt_tgsigqueueinfo _x32_compat_sys_rt_tgsigqueueinfo
378 537 x32 recvmmsg _x32_compat_sys_recvmmsg
379 538 x32 sendmmsg _x32_compat_sys_sendmmsg
380 539 x32 process_vm_readv _x32_compat_sys_process_vm_readv
381 540 x32 process_vm_writev _x32_compat_sys_process_vm_writev
382 541 x32 setsockopt _x32_compat_sys_setsockopt
383 542 x32 getsockopt _x32_compat_sys_getsockopt
384 543 x32 io_setup _x32_compat_sys_io_setup
385 544 x32 io_submit _x32_compat_sys_io_submit
386 545 x32 execveat _x32_compat_sys_execveat/ptregs
387 546 x32 preadv2 _x32_compat_sys_preadv64v2
388 547 x32 pwritev2 _x32_compat_sys_pwritev64v2
389 548 64 hello sys hello
390 549 64 shlokhello shlokhello
```

In kernel all system calls are identified by a unique number.

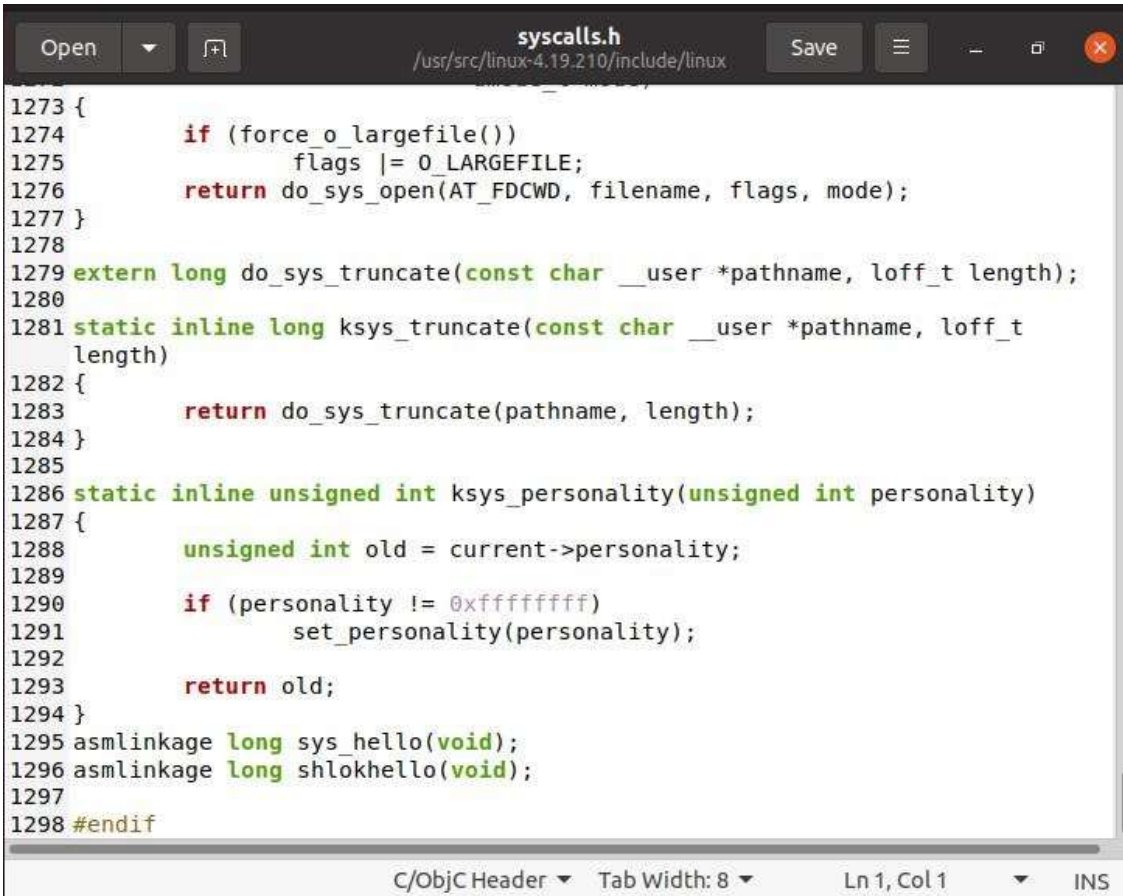
Here we assigned a unique number to our own system function call i.e., **549**.

Step 7) Add new system call to the system call header file

```
cd /usr/src/linux-4.19.210/
```

```
cd include/linux/
```

```
gedit syscalls.h
```



```
1273 {
1274     if (force_o_largefile())
1275         flags |= O_LARGEFILE;
1276     return do_sys_open(AT_FDCWD, filename, flags, mode);
1277 }
1278
1279 extern long do_sys_truncate(const char __user *pathname, loff_t length);
1280
1281 static inline long ksys_truncate(const char __user *pathname, loff_t
length)
1282 {
1283     return do_sys_truncate(pathname, length);
1284 }
1285
1286 static inline unsigned int ksys_personality(unsigned int personality)
1287 {
1288     unsigned int old = current->personality;
1289
1290     if (personality != 0xffffffff)
1291         set_personality(personality);
1292
1293     return old;
1294 }
1295 asmlinkage long sys_hello(void);
1296 asmlinkage long shlokhello(void);
1297
1298 #endif
```

This defines the prototype of the function of our system call.

Step 8) Compile the Kernel

```
cd /usr/src/linux-4.19.210/
```

```
make menuconfig
```

This command is used to configure the Linux kernel. A pop up will come just check that ext4 in file system is included.

```
make -j4
```

This command will start compiling the kernel. Now j4 in **make -j4** makes use of 4 cores in system to make this file. The more cores we use the faster the compilation will take place. -jn tells number of cores to be used and n cannot be greater than number of cores in your machine

Step 9) Update or Install the kernel

```
make modules_install install
```

The above command will make changes in the kernel according to the files which we have compiled above, and the changes will be updated in the 4 files in **/boot** directory: -

1. System.map-4.17.4
2. vmlinuz-4.17.4
3. initrd.img-4.17.4
4. config-4.17.4

Now reboot the system to see all the changes in kernel.

Step 10) Test the system call

Code:

```
1 #include<stdio.h>
2 #include<linux/kernel.h>
3 #include<sys/syscall.h>
4 #include<unistd.h>
5 int main()
6 {
7     long int no=0;
8     no=syscall(549);
9     printf("Ans is %d:",no);
10    return 0;
11 }
```

Output:

```
shlok@shlok:~$ ./q1
Ans is 0:shlok@shlok:~$ dmesg
```

```
shlok@shlok: ~
[ 6.589184] systemd[1]: Condition check resulted in Rebuild Hardware Database
being skipped.
[ 6.589338] systemd[1]: Condition check resulted in Platform Persistent Stora
ge Archival being skipped.
[ 6.591693] systemd[1]: Starting Load/Save Random Seed...
[ 6.597328] systemd[1]: Starting Create System Users...
[ 6.599947] systemd[1]: Finished Uncomplicated firewall.
[ 6.755779] systemd[1]: Started Journal Service.
[ 6.835840] systemd-journald[1173]: Received client request to flush runtime
journal.
[ 6.892784] Adding 703976k swap on /swapfile. Priority:-2 extents:3 across:7
20360k
[ 11.934040] ethtool (1248) used greatest stack depth: 13456 bytes left
[ 15.968600] e2scrub_all (1342) used greatest stack depth: 13248 bytes left
[ 24.783637] IPv6: ADDRCONF(NETDEV_UP): enp0s3: link is not ready
[ 24.789655] e1000: enp0s3 NIC Link is Up 1000 Mbps Full Duplex, Flow Control:
RX
[ 24.791592] IPv6: ADDRCONF(NETDEV_UP): enp0s3: link is not ready
[ 24.791609] IPv6: ADDRCONF(NETDEV_CHANGE): enp0s3: link becomes ready
[ 70.412982] rfkill: input handler disabled
[ 2757.974504] kworker/dying (319) used greatest stack depth: 13032 bytes left
[ 10752.982342] kworker/dying (2990) used greatest stack depth: 12752 bytes left
[ 11735.885059] Welcome to Linux shlok
```

Question 2) Write syscall which will receive string parameter and print it along with some message to kernel logs

Step 1) Go to the kernel directory where we want to make a new system call

```
cd /usr/src/linux-4.19.210/
```

With this we change our and work in the kernel directory.

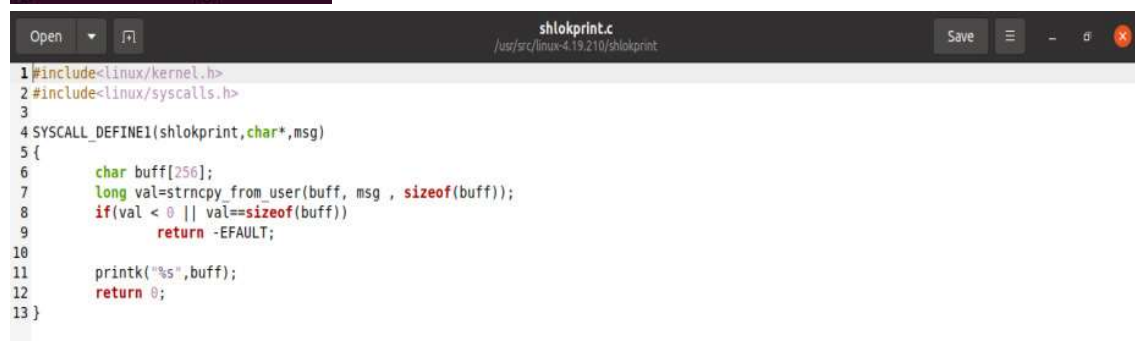
Step 2) Make the directory in which you will write the source code of the system call and go into that directory.

```
mkdir shlokprint
```

```
cd shlokprint
```

Step 3) Create shlokprint.c in this directory.

```
gedit shlokprint.c
```



```
1 #include<linux/kernel.h>
2 #include<linux/syscalls.h>
3
4 SYSCALL_DEFINE1(shlokprint,char*,msg)
5 {
6     char buff[256];
7     long val=strncpy_from_user(buff, msg , sizeof(buff));
8     if(val < 0 || val==sizeof(buff))
9         return -EFAULT;
10
11     printk("%s",buff);
12     return 0;
13 }
```

- **SYSCALL_DEFINE1** is used to define function with parameters. Here **1** means the function has only 1 parameter. The first parameter is the function name and it makes function name as **sys_shlokprint(char *msg)**.
- The parameters are defines as pairs in SYSCALL_DEFINE(). Now (char* ,msg) is same as char *msg in normal ways.
- When data is passed as a pointer then pointer address is send relative to user space. So when we directly try to access that address in kernel space it gives us error.
- To avoid this error we are taking help of **strncpy_from_user** looks for message pointer in user space and copy the content to the buffer variable of the kernel space .
- **val** tells us how much data is copied in the buffer. If val<0 then there is error in copy and when val=sizeof(buffer) then there is overflow so we have sent EFAULT when this occurs
- **printk** is a function like printf but it is used to print in kernel log.

Step 4) Create Makefile

```
gedit Makefile
```



```
1 obj-y := shlokprint.o
```

This is to ensure that shlokprint.c is compiled when the kernel is compiled.

Step 5) Edit the Makefile of Kernel

```
cd /usr/src/linux-4.19.210/
```

```
nano Makefile
```

Look for the line which looks like below

core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/

In this line add shlokprint/

```
core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ hello/ shlokhello/ shlokprint/
```

- By adding shlokprint/ in this line we are telling the compiler of the kernel that our new system call is defined in shlokprint directory.
- It will tell the compiler to make the Makefile in this directory which was mentioned above.

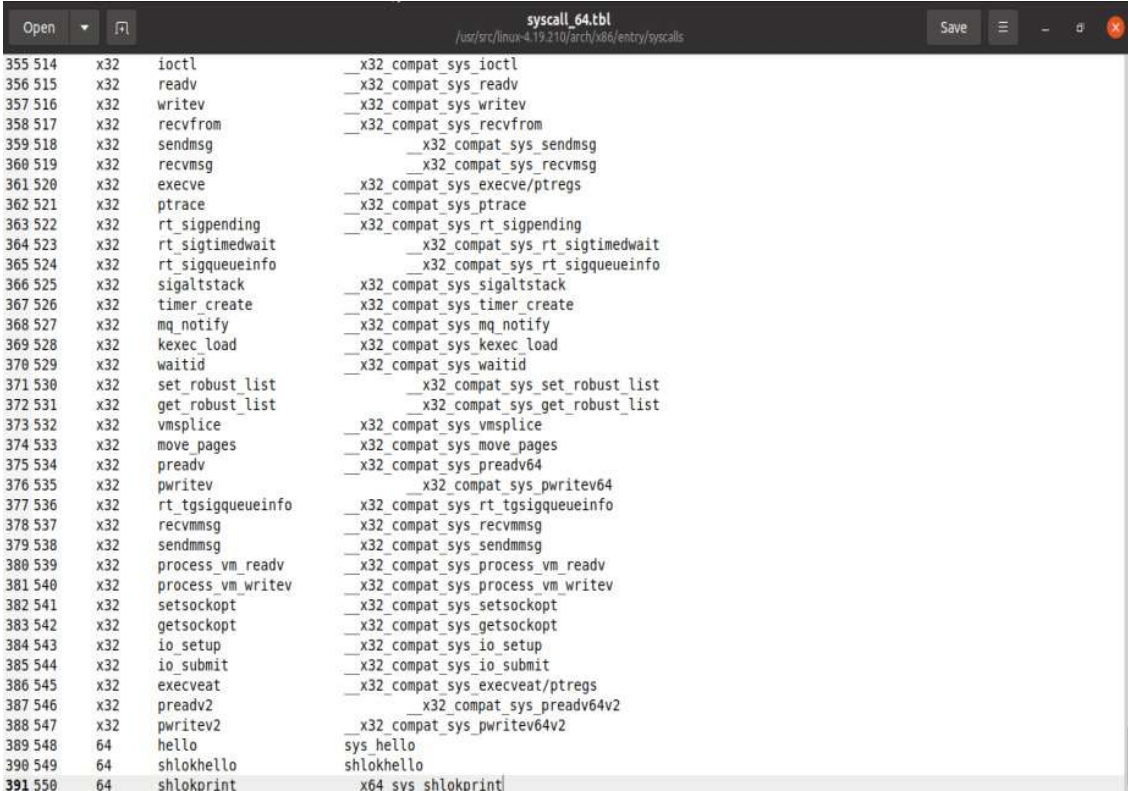
Step 6) Add the new system call to the system call table of kernel

```
cd arch/x86/entry/syscalls/
```

The system call table of the kernel lies in above directory. Change your current directory with the above command.

As my system is 64 bit so add entry in syscall_64.tbl .

```
gedit syscall_64.tbl
```



syscall_64.tbl	
/usr/src/linux-4.19.210/arch/x86/entry/syscalls	
355 514	x32 ioctl _x32_compat_sys_ioctl
356 515	x32 readv _x32_compat_sys_readv
357 516	x32 writev _x32_compat_sys_writev
358 517	x32 recvfrom _x32_compat_sys_recvfrom
359 518	x32 sendmsg _x32_compat_sys_sendmsg
360 519	x32 recvmsg _x32_compat_sys_recvmsg
361 520	x32 execve _x32_compat_sys_execve/ptregs
362 521	x32 ptrace _x32_compat_sys_ptrace
363 522	x32 rt_sigpending _x32_compat_sys_rt_sigpending
364 523	x32 rt_sigtimedwait _x32_compat_sys_rt_sigtimedwait
365 524	x32 rt_sigqueueinfo _x32_compat_sys_rt_sigqueueinfo
366 525	x32 sigaltstack _x32_compat_sys_sigaltstack
367 526	x32 timer_create _x32_compat_sys_timer_create
368 527	x32 mq_notify _x32_compat_sys_mq_notify
369 528	x32 kexec_load _x32_compat_sys_kexec_load
370 529	x32 waitid _x32_compat_sys_waitid
371 530	x32 set_robust_list _x32_compat_sys_set_robust_list
372 531	x32 get_robust_list _x32_compat_sys_get_robust_list
373 532	x32 vmsplice _x32_compat_sys_vmsplice
374 533	x32 move_pages _x32_compat_sys_move_pages
375 534	x32 preadv _x32_compat_sys_preadv64
376 535	x32 pwritev _x32_compat_sys_pwritev64
377 536	x32 rt_tsigqueueinfo _x32_compat_sys_rt_tsigqueueinfo
378 537	x32 recvmsg _x32_compat_sys_recvmsg
379 538	x32 sendmsg _x32_compat_sys_sendmsg
380 539	x32 process_vm_readv _x32_compat_sys_process_vm_readv
381 540	x32 process_vm_writev _x32_compat_sys_process_vm_writev
382 541	x32 setsockopt _x32_compat_sys_setsockopt
383 542	x32 getsockopt _x32_compat_sys_getsockopt
384 543	x32 io_setup _x32_compat_sys_io_setup
385 544	x32 io_submit _x32_compat_sys_io_submit
386 545	x32 execveat _x32_compat_sys_execveat/ptregs
387 546	x32 preadv2 _x32_compat_sys_preadv64v2
388 547	x32 pwritev2 _x32_compat_sys_pwritev64v2
389 548	64 hello sys_hello
390 549	64 shlokhello shlokhello
391 550	64 shlokprint x64_sys_shlokprint

In kernel all system calls are identified by a unique number.

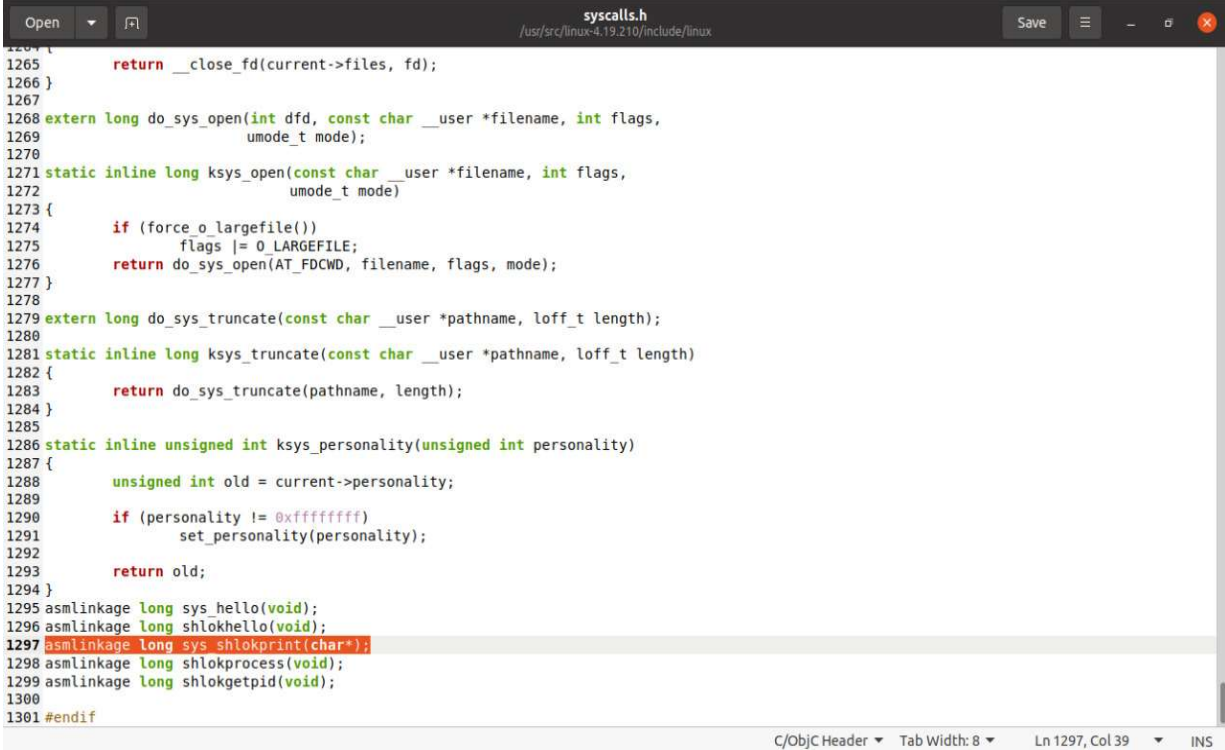
Here we assigned a unique number to our own system function call i.e., 550.

Step 7) Add new system call to the system call header file

```
cd /usr/src/linux-4.19.210/
```

```
cd include/linux/
```

```
gedit syscalls.h
```



```
1265     return __close_fd(current->files, fd);
1266 }
1267
1268 extern long do_sys_open(int dfd, const char __user *filename, int flags,
1269                        umode_t mode);
1270
1271 static inline long ksys_open(const char __user *filename, int flags,
1272                             umode_t mode)
1273 {
1274     if (force_o_largefile())
1275         flags |= O_LARGEFILE;
1276     return do_sys_open(AT_FDCWD, filename, flags, mode);
1277 }
1278
1279 extern long do_sys_truncate(const char __user *pathname, loff_t length);
1280
1281 static inline long ksys_truncate(const char __user *pathname, loff_t length)
1282 {
1283     return do_sys_truncate(pathname, length);
1284 }
1285
1286 static inline unsigned int ksys_personality(unsigned int personality)
1287 {
1288     unsigned int old = current->personality;
1289
1290     if (personality != 0xffffffff)
1291         set_personality(personality);
1292
1293     return old;
1294 }
1295 asmlinkage long sys_hello(void);
1296 asmlinkage long shlokhello(void);
1297 asmlinkage long sys_shlokprint(char*);
1298 asmlinkage long shlokprocess(void);
1299 asmlinkage long shlokgetpid(void);
1300
1301 #endif
```

This defines the prototype of the function of our system call.

Step 8) Compile the Kernel

```
cd /usr/src/linux-4.19.210/
```

```
make menuconfig
```

This command is used to configure the Linux kernel. A pop up will come just check that ext4 in file system is included.

```
make -j4
```

This command will start compiling the kernel. Now j4 in **make -j4** makes use of 4 cores in system to make this file. The more cores we use the faster the compilation will take place. -jn tells number of cores to be used and n cannot be greater than number of cores in your machine

Step 9) Update or Install the kernel

```
make modules_install install
```

The above command will make changes in the kernel according to the files which we have compiled above, and the changes will be updated in the 4 files in **/boot** directory: -

2. System.map-4.17.4
3. vmlinuz-4.17.4
4. initrd.img-4.17.4
5. config-4.17.4

Now reboot the system to see all the changes in kernel.

Step 10) Test the system call

Code:

```
Q2.c
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <sys/syscall.h>
4
5
6
7 int main(int argc, char **argv)
8 {
9     char str[20] = "Omkar";
10    // Make the greet_o system call
11    int ret_val = syscall(550, str);
12    if (ret_val < 0)
13    {
14        printf("[ERR] Ret val: %d\n", ret_val);
15        return 1;
16    }
17    printf("greet_o system call message: %d\n", ret_val);
18    return 0;
19 }
```

Output:

```
shlok@shlok:~$ gcc Q2.c -o q2
shlok@shlok:~$ ./q2
greet_o system call message: 0

shlok@shlok:~$ gcc Q2.c -o q2
shlok@shlok:~$ ./q2
greet_o system call message: 0
RX
[ 24.791592] IPv6: ADDRCONF(NETDEV_UP): enp0s3: link is not ready
[ 24.791609] IPv6: ADDRCONF(NETDEV_CHANGE): enp0s3: link becomes ready
[ 70.412982] rfkill: input handler disabled
[ 2757.974504] kworker/dying (319) used greatest stack depth: 13032 bytes left
[10752.982342] kworker/dying (2990) used greatest stack depth: 12752 bytes left
[11735.885059] Welcome to Linux shlok
[11758.476060] Welcome to Linux shlok
[11863.018553] Omkar
```

Question 3) Write system call to print the parent process id and current process id upon calling it

Step 1) Go to the kernel directory where we want to make a new system call

```
cd /usr/src/linux-4.19.210/
```

With this we change our and work in the kernel directory.


Step 2) Make the directory in which you will write the source code of the system call and go into that directory.

```
mkdir shlokprocess
```

```
cd shlokprocess
```

Step 3) Create shlokprocess.c in this directory.

```
gedit shlokprocess.c
```

A screenshot of a text editor window titled 'shlokprocess.c' with the path '/usr/src/linux-4.19.210/shlokprocess'. The code is as follows:

```
1
2 #include <linux/kernel.h>
3
4 #include <linux/sched.h>
5
6
7 asmlinkage long shlokprocess(void)
8 {
9     printk("CID:%d PID:%d", current->pid, task_ppid_nr(current));
10    //printk("Parent Process id:%d", task_ppid_nr(current));
11    return 0;
12 }
```

- **asmlinkage** is a keyword which is used to indicate that all parameters of the function would be available of stack.
- **void** indicates function that shlokprocess() does not take any parameters.
- **printk** is a function like printf but it is used to print in kernel log.
- **current** is global pointer variable which of **struct task_struct type**. It points to the current process' data and with the help of pointer we are easily accessing the process' pid which is available in it's task_struct data and return its value.
- **task_ppid_nr** helps us to get parent id of the pointer which we pass to it.

Step 4) Create Makefile

```
gedit Makefile
```

A screenshot of a text editor window titled 'Makefile' with the path '/usr/src/linux-4.19.210/shlokprocess'. The code is as follows:

```
1 obj-y := shlokprocess.o
```

This is to ensure that shlokprocess.c is compiled when the kernel is compiled.

Step 5) Edit the Makefile of Kernel

```
cd /usr/src/linux-4.19.210/
```

```
nano Makefile
```

Look for the line which looks like below

```
core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/
```

In this line add shlokprocess/

```
core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ hello/ shlokhell/ shlokprint/ shlokprocess/
```

- By adding shlokprocess/ in this line we are telling the compiler of the kernel that our new system call is defined in shlokprocess directory.
- It will tell the compiler to make the Makefile in this directory which was mentioned above.

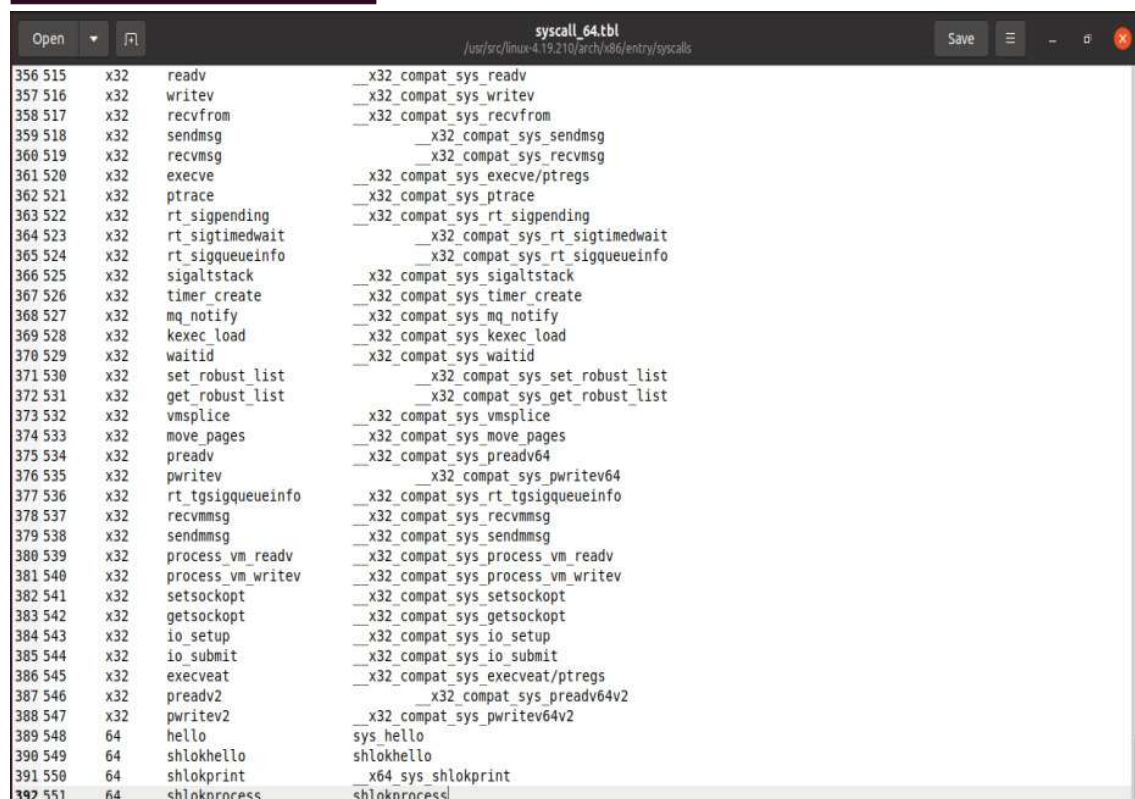
Step 6) Add the new system call to the system call table of kernel

```
cd arch/x86/entry/syscalls/
```

The system call table of the kernel lies in above directory. Change your current directory with the above command.

As my system is 64 bit so add entry in syscall_64.tbl .

```
gedit syscall_64.tbl
```



Line	Unique Number	System Call	Kernel Function
356	515	readv	__x32_compat_sys_readv
357	516	writew	__x32_compat_sys_writew
358	517	recvfrom	__x32_compat_sys_recvfrom
359	518	sendmsg	__x32_compat_sys_sendmsg
360	519	recvmsg	__x32_compat_sys_recvmsg
361	520	execve	__x32_compat_sys_execve/ptregs
362	521	ptrace	__x32_compat_sys_ptrace
363	522	rt_sigpending	__x32_compat_sys_rt_sigpending
364	523	rt_sigtimedwait	__x32_compat_sys_rt_sigtimedwait
365	524	rt_sigqueueinfo	__x32_compat_sys_rt_sigqueueinfo
366	525	sigaltstack	__x32_compat_sys_sigaltstack
367	526	timer_create	__x32_compat_sys_timer_create
368	527	mq_notify	__x32_compat_sys_mq_notify
369	528	kexec_load	__x32_compat_sys_kexec_load
370	529	waitid	__x32_compat_sys_waitid
371	530	set_robust_list	__x32_compat_sys_set_robust_list
372	531	get_robust_list	__x32_compat_sys_get_robust_list
373	532	vmsplice	__x32_compat_sys_vmsplice
374	533	move_pages	__x32_compat_sys_move_pages
375	534	preadv	__x32_compat_sys_preadv64
376	535	pwritev	__x32_compat_sys_pwritev64
377	536	rt_tgsigqueueinfo	__x32_compat_sys_rt_tgsigqueueinfo
378	537	recvmsg	__x32_compat_sys_recvmsg
379	538	sendmsg	__x32_compat_sys_sendmsg
380	539	process_vm_readv	__x32_compat_sys_process_vm_readv
381	540	process_vm_writev	__x32_compat_sys_process_vm_writev
382	541	setsockopt	__x32_compat_sys_setsockopt
383	542	getsockopt	__x32_compat_sys_getsockopt
384	543	io_setup	__x32_compat_sys_io_setup
385	544	io_submit	__x32_compat_sys_io_submit
386	545	execveat	__x32_compat_sys_execveat/ptregs
387	546	preadv2	__x32_compat_sys_preadv64v2
388	547	pwritev2	__x32_compat_sys_pwritev64v2
389	548	hello	sys_hello
390	549	shlokhell	shlokhell
391	550	shlokprint	__x64_sys_shlokprint
392	551	shlokprocess	shlokprocess

In kernel all system calls are identified by a unique number.

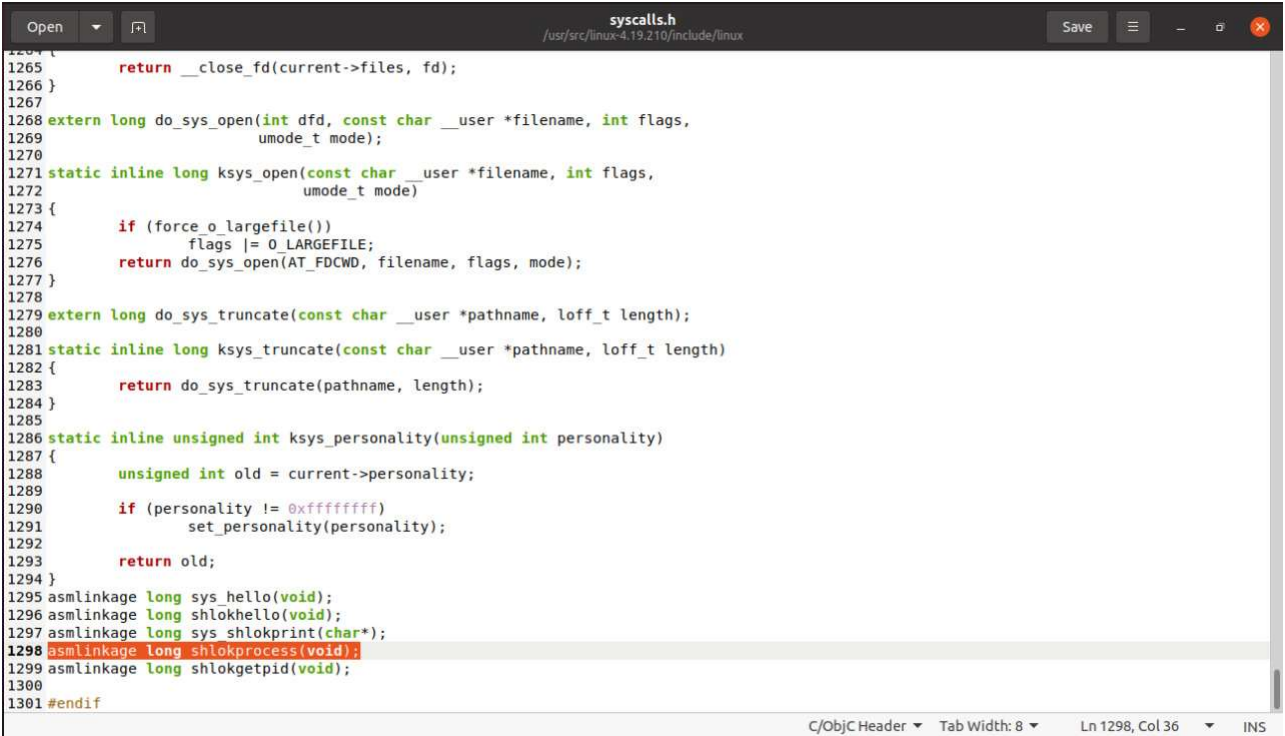
Here we assigned a unique number to our own system function call i.e., **551**.

Step 7) Add new system call to the system call header file

```
cd /usr/src/linux-4.19.210/
```

```
cd include/linux/
```

```
gedit syscalls.h
```



```
1264 {
1265     return __close_fd(current->files, fd);
1266 }
1267
1268 extern long do_sys_open(int dfd, const char __user *filename, int flags,
1269                        umode_t mode);
1270
1271 static inline long ksys_open(const char __user *filename, int flags,
1272                             umode_t mode)
1273 {
1274     if (force_o_largefile())
1275         flags |= O_LARGEFILE;
1276     return do_sys_open(AT_FDCWD, filename, flags, mode);
1277 }
1278
1279 extern long do_sys_truncate(const char __user *pathname, loff_t length);
1280
1281 static inline long ksys_truncate(const char __user *pathname, loff_t length)
1282 {
1283     return do_sys_truncate(pathname, length);
1284 }
1285
1286 static inline unsigned int ksys_personality(unsigned int personality)
1287 {
1288     unsigned int old = current->personality;
1289
1290     if (personality != 0xffffffff)
1291         set_personality(personality);
1292
1293     return old;
1294 }
1295 asmlinkage long sys_hello(void);
1296 asmlinkage long shlokhello(void);
1297 asmlinkage long sys_shlokprint(char*);
1298 asmlinkage long shlokprocess(void);
1299 asmlinkage long shlokgetpid(void);
1300
1301 #endif
```

This defines the prototype of the function of our system call.

Step 8) Compile the Kernel

```
cd /usr/src/linux-4.19.210/
```

```
make menuconfig
```

This command is used to configure the Linux kernel. A pop up will come just check that ext4 in file system is included.

```
make -j4
```

This command will start compiling the kernel. Now j4 in **make -j4** makes use of 4 cores in system to make this file. The more cores we use the faster the compilation will take place. -jn tells number of cores to be used and n cannot be greater than number of cores in your machine

Step 9) Update or Install the kernel

```
make modules_install install
```

The above command will make changes in the kernel according to the files which we have compiled above, and the changes will be updated in the 4 files in **/boot** directory: -

3. System.map-4.17.4
4. vmlinuz-4.17.4
5. initrd.img-4.17.4
6. config-4.17.4

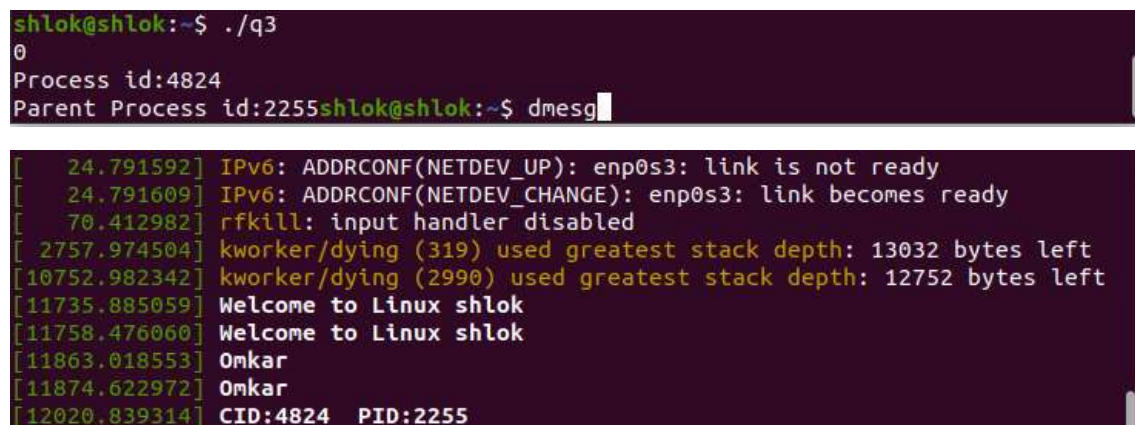
Now reboot the system to see all the changes in kernel.

Code:



```
1#include <stdio.h>
2#include <sys/types.h>
3#include <unistd.h>
4#include <sys/syscall.h>
5
6int main()
7{
8    pid_t process_id=getpid();
9    pid_t parent_pro_id=getppid();
10    long int num=syscall(551);
11    printf("%ld\n",num);
12    printf("Process id:%d",process_id);
13    printf("\nParent Process id:%d",parent_pro_id);
14    return 0;
15}
```

Output:



```
shlok@shlok:~$ ./q3
0
Process id:4824
Parent Process id:2255shlok@shlok:~$ dmesg
[ 24.791592] IPv6: ADDRCONF(NETDEV_UP): enp0s3: link is not ready
[ 24.791609] IPv6: ADDRCONF(NETDEV_CHANGE): enp0s3: link becomes ready
[ 70.412982] rfkill: input handler disabled
[ 2757.974504] kworker/dying (319) used greatest stack depth: 13032 bytes left
[10752.982342] kworker/dying (2990) used greatest stack depth: 12752 bytes left
[11735.885059] Welcome to Linux shlok
[11758.476060] Welcome to Linux shlok
[11863.018553] Onkar
[11874.622972] Onkar
[12020.839314] CID:4824 PID:2255
```

Q) Are both process ids i.e. is current process id and parent process id same or different? Why? What are your observations.

Ans:

Every process has a unique process id which is assigned by OS during the process creation. Now as we know that when a function call takes place during execution of process the executing code changes but the running process remains same. Similarly when a system call is called no new process is created so the currently running process does not changes so Process ID also not changes and the system call is part of this process.

As you can see that in above example the system call shlokprocess() was called by q3. The process id of q3 and the CID (Current ID) of system call is same i.e. 4824. Therefore parent of q3 is same as parent of system call as systemcall is part of that is why the PID (Parent id) printed in system call is same as the parent of q3 i.e. 2255.

Question 4) Write system call to execute some predefined system call from your written system call.

Step 1) Go to the kernel directory where we want to make a new system call

```
cd /usr/src/linux-4.19.210/
```

With this we change our and work in the kernel directory.

Step 2) Make the directory in which you will write the source code of the system call and go into that directory.

```
mkdir shlokgetpid
```

```
cd shlokgetpid
```

Step 3) Create shlokgetpid.c in this directory.

```
gedit shlokgetpid.c
```



```
shlokgetpid.c
/usr/src/linux-4.19.210/shlokgetpid

1 #include <linux/module.h>
2 #include <linux/kernel.h>
3 #include <linux/unistd.h>
4 #include <linux/syscalls.h>
5 #include <linux/fs.h>
6 #include <linux/sched.h>
7
8
9 asmlinkage long shlokgetpid(void)
10 {
11     return current->pid;
12 }
```

- **asmlinkage** is a keyword which is used to indicate that all parameters of the function would be available of stack.
- **void** indicates function that shlokgetpid() does not take any parameters.
- **current** is global pointer variable which of **struct task_struct type**. It points to the current process' data and with the help of pointer we are easily accessing the process' pid which is available in it's task_struct data and return its value.

Step 4) Create Makefile

```
gedit Makefile
```



```
Makefile
/usr/src/linux-4.19.210/shlokgetpid

1 obj-y := shlokgetpid.o
```

This is to ensure that shlokgetpid.c is compiled when the kernel is compiled.

Step 5) Edit the Makefile of Kernel

```
cd /usr/src/linux-4.19.210/
```

```
nano Makefile
```

Look for the line which looks like below

```
core-y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/
```

In this line add shlokgetpid/

```
core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ hello/ shlokhello/ shlokprint/ shlokprocess/ shlokgetpid/
```

- By adding shlokgetpid/ in this line we are telling the compiler of the kernel that our new system call is defined in shlokgetpid directory.
- It will tell the compiler to make the Makefile in this directory which was mentioned above.

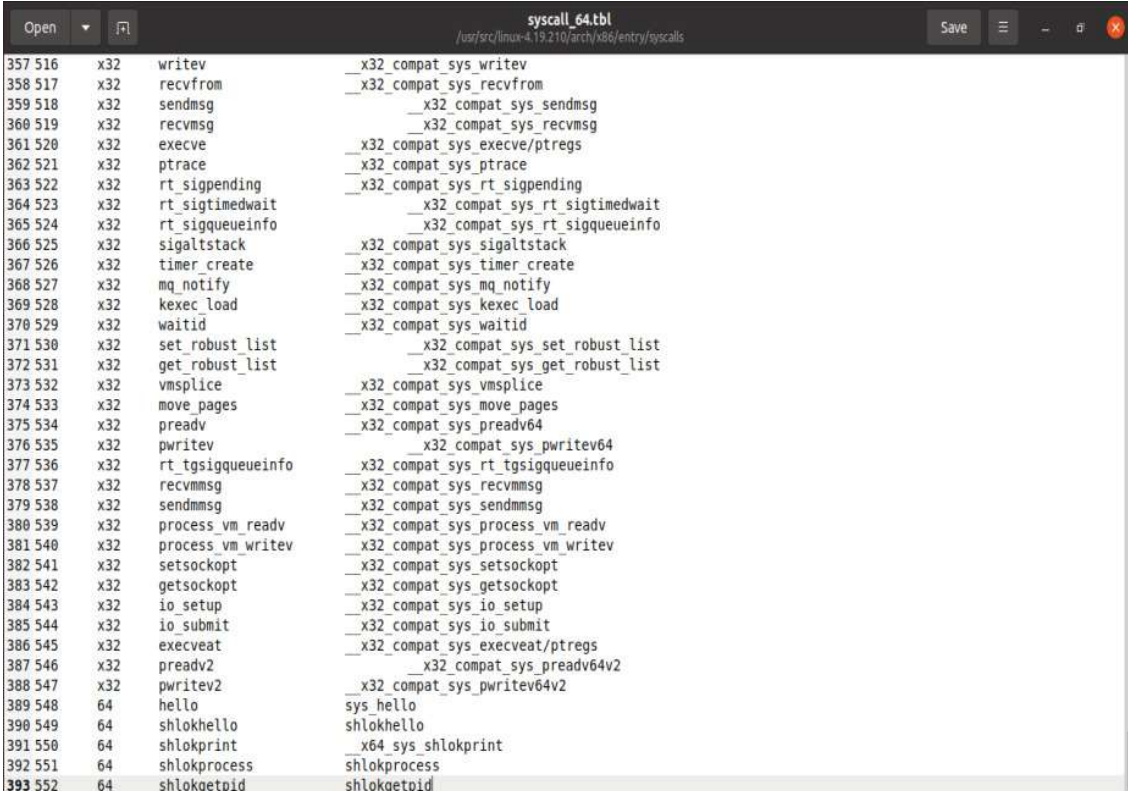
Step 6) Add the new system call to the system call table of kernel

```
cd arch/x86/entry/syscalls/
```

The system call table of the kernel lies in above directory. Change your current directory with the above command.

As my system is 64 bit so add entry in syscall_64.tbl .

```
gedit syscall_64.tbl
```



syscall_64.tbl			
/usr/src/linux-4.19.210/arch/x86/entry/syscalls			
357	516	x32	writev
358	517	x32	recvfrom
359	518	x32	sendmsg
360	519	x32	recvmsg
361	520	x32	execve
362	521	x32	ptrace
363	522	x32	rt_sigpending
364	523	x32	rt_sigtimedwait
365	524	x32	rt_sigqueueinfo
366	525	x32	sigaltstack
367	526	x32	timer_create
368	527	x32	mq_notify
369	528	x32	kexec_load
370	529	x32	waitid
371	530	x32	set_robust_list
372	531	x32	get_robust_list
373	532	x32	vmsplice
374	533	x32	move_pages
375	534	x32	preadv
376	535	x32	pwritev
377	536	x32	rt_tgsigqueueinfo
378	537	x32	recvmsg
379	538	x32	sendmsg
380	539	x32	process_vm_readv
381	540	x32	process_vm_writev
382	541	x32	setsockopt
383	542	x32	getsockopt
384	543	x32	io_setup
385	544	x32	io_submit
386	545	x32	execveat
387	546	x32	preadv2
388	547	x32	pwritev2
389	548	64	hello
390	549	64	shlokhello
391	550	64	shlokprint
392	551	64	shlokprocess
393	552	64	shlokgetpid

In kernel all system calls are identified by a unique number.


Here we assigned a unique number to our own system function call i.e., 552.

Step 7) Add new system call to the system call header file

```
cd /usr/src/linux-4.19.210/
```

```
cd include/linux/
```

```
gedit syscalls.h
```



```
1265     return __close_fd(current->files, fd);
1266 }
1267
1268 extern long do_sys_open(int dfd, const char __user *filename, int flags,
1269                        umode_t mode);
1270
1271 static inline long ksys_open(const char __user *filename, int flags,
1272                             umode_t mode)
1273 {
1274     if (force_o_largefile())
1275         flags |= O_LARGEFILE;
1276     return do_sys_open(AT_FDCWD, filename, flags, mode);
1277 }
1278
1279 extern long do_sys_truncate(const char __user *pathname, loff_t length);
1280
1281 static inline long ksys_truncate(const char __user *pathname, loff_t length)
1282 {
1283     return do_sys_truncate(pathname, length);
1284 }
1285
1286 static inline unsigned int ksys_personality(unsigned int personality)
1287 {
1288     unsigned int old = current->personality;
1289
1290     if (personality != 0xffffffff)
1291         set_personality(personality);
1292
1293     return old;
1294 }
1295 asmlinkage long sys_hello(void);
1296 asmlinkage long shlokhello(void);
1297 asmlinkage long sys_shlokprint(char*);
1298 asmlinkage long shlokprocess(void);
1299 asmlinkage long shlokgetpid(void);
1300
1301 #endif
```

This defines the prototype of the function of our system call.

Step 8) Compile the Kernel

```
cd /usr/src/linux-4.19.210/
```

```
make menuconfig
```

This command is used to configure the Linux kernel. A pop up will come just check that ext4 in file system is included.

```
make -j4
```

This command will start compiling the kernel. Now j4 in **make -j4** makes use of 4 cores in system to make this file. The more cores we use the faster the compilation will take place. -jn tells number of cores to be used and n cannot be greater than number of cores in your machine

Step 9) Update or Install the kernel

```
make modules_install install
```

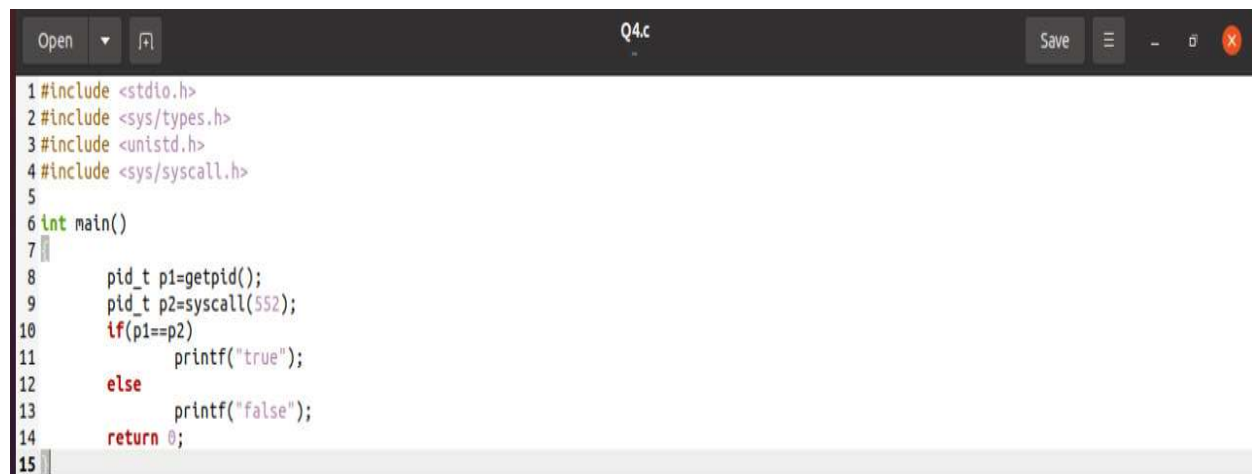
The above command will make changes in the kernel according to the files which we have compiled above, and the changes will be updated in the 4 files in **/boot** directory: -

4. System.map-4.17.4
5. vmlinuz-4.17.4
6. initrd.img-4.17.4
7. config-4.17.4

Now reboot the system to see all the changes in kernel.

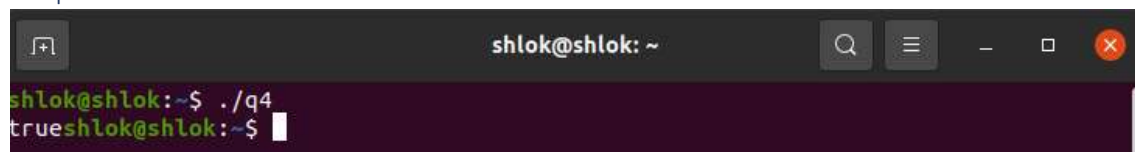
Step 10) Test the system call

Code:



```
1 #include <stdio.h>
2 #include <sys/types.h>
3 #include <unistd.h>
4 #include <sys/syscall.h>
5
6 int main()
7 {
8     pid_t p1=getpid();
9     pid_t p2=syscall(552);
10    if(p1==p2)
11        printf("true");
12    else
13        printf("false");
14    return 0;
15 }
```

Output:



```
shlok@shlok: ~$ ./q4
true
shlok@shlok: ~$
```